

COMP90024 Cluster and Cloud Computing

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

University of Melbourne, March 22, 2018

lev.lafayette@unimelb.edu.au

"This is an advanced course but we get mixed bag: students that have 5+ years of MPI programming on supercomputers, to students that have only done Java on Windows."

- Some background on supercomputing, high performance computing, parallel computing, scientific computing (there is overlap, but they're not the same thing).
- An introduction to Spartan, Cloud hybrid system
- Logging in, help, and enviro
- Job submission with Slurm workload manager; sions, multicore, job arrays, job dependencies, interactive jobs.
- Parallel programming with shared memory and threads (OpenMP) and distributed memory and message passing (OpenMPI)
- Tantalising hints about more advanced material on message passing routines.

Assignment, Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

'Supercomputer' arbitrary term with no specific definition. In general use it means any single computer system (itself a contested term) that has exceptional processing power for its time. A well-adopted metric is the number of floating-point operations per second (FLOPS) such a system can carry out.

Supercomputers, like any other computing system, have improved significantly over time. The Top500 list is based on FLOPS using LINPACK. HPC Challenge is a broader, more interesting metric. The current number #1 system is Sunway TaihuLight, a supercomputer operated by China's National Super Comput

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

1994: 170.40 GFLOPS
1996: 368.20 GFLOPS
1997: 1.338 TFLOPS
1999: 2.3796 TFLOPS
2000: 7.226 TFLOPS
2004: 70.72 TFLOPS
2005: 280.6 TFLOPS
2007: 478.2 TFLOPS
2008: 1.105 PFLOPS
2009: 1.759 PFLOPS
2010: 2.566 PFLOPS
2011: 10.51 PFLOPS
2012: 17.59 PFLOPS
2013: 33.86 PFLOPS
2014: 33.86 PFLOPS
2015: 33.86 PFLOPS
2016: 93.01 PFLOPS
2017: 93.01 PFLOPS (125.46 PFLOPS peak)

High-performance computing (HPC) is any computer system whose architecture allows for above average performance. A system that is one of the most powerful in the world, but is poorly designed, could be a "supercomputer".

Clustered computing is when two or more computers serve a single resource. This improves performance and provides redundancy in case of failure system. To describe simply, there are a collection of smaller computers strapped together with a high-speed local network (e.g., Myrinet, InfiniBand, 10 Gigabit Ethernet) system could certainly be used. Even a cluster of Raspberry Pi w Lego chassis (University of Southampton, 2012)!

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Horse and cart as a computer system and the load as the computing tasks. Efficient arrangement, bigger horse and cart, or a teamster? The clustered HPC is the most efficient, economical, and scalable method, and for that reason it dominates supercomputing today.

With a cluster architecture, applications can be more easily parallelised across them. Parallel computing refers to the submission of jobs or processes over multiple processors and by splitting up the data or tasks between them (random number generation as data parallel, driving a vehicle as task parallel).

Research computing is the software applications used by a research community to aid research. Does not necessarily equate with high performance computing, or the use of clusters. This skills gap is a major problem and must be addressed because as the volume, velocity, and variety of datasets increases then research


Assignment Project Exam Help

<https://eduassistpro.github.io/>

Computational capacity does have a priority (the system prior to use), in order for that capacity to be realised in terms of usage a skill set can be developed. The core issue is that high performance compute clusters is just speed and power but also usage, productivity, correctness, and reproducibility.

(image from Lawrence Livermore National Laboratory)

There is nascent research that shows a strong correlation between research output and availability of HPC facilities. (Apon et al 2010)



Researchers from Monash University, the Peter MacCallum Cancer Institute in Melbourne, the Birkbeck College in London, and VPAC in 2010 unravelled the structure the protein perforin to determine how pathogenic cells are attacked by white blood cells.

Assignment Project Exam Help

In 2015 researchers from VLSCI antifreeze proteins bind to ice to prevent it growing has important implications for extending climate or and protecting crops from frost damage.

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](#)

In 2016 CSIRO researchers successfully manipulated the behaviour of Metallic Organic Frameworks to control their structure and alignment which provides opportunities for real-time and implantable medical electric devices.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

In November 2017 of the Top 500 Supercomputers worldwide,
every single machine used Linux.

<http://www.zdnet.com/article/linux-totally-dominates-supercomputers/>

The command-line interface provides a great deal more power
and is very resource efficient.

Assignment Project Exam Help

GNU/Linux scales and does so

Critical software such as the Me
and nearly all scientific progra

GNU/Linux.

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The operating system and many applications are provided
as "free and open source", which means that not only are
there are some financial savings, were also much better
placed to improve, optimize and maintain specific programs.

Free or open source software (not always the same thing)
can be can be compiled from source for the specific
hardware and operating system configuration, and can be
optimised according to compiler flags. There is necessary
where every clock cycle is important.

It is possible to illustrate the degree of parallelisation by using Flynn's Taxonomy of Computer Systems (1966), where each process is considered as the execution of a pool of instructions (instruction stream) on a pool of data (data stream).

Over time computing systems have moved towards multi-processor, multi-core, and often multi-threaded and multi-node systems.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

The engineering imperative to t comes down to heat. From the mid-2000s clock speed on CPUs have largely stalled.

Add WeChat edu_assist_pro

Some trends include GPGPU development, massive multicore systems (e.g., The Angstrom Project, the Tile CPU with 1000 cores) and massive network connectivity and shared resources (e.g., Plan9 Operating System).

(Image from Dr. Mark Meyer, Canisius College)

Parallel programming and multicore systems should mean better performance. This can be expressed a ratio called speedup

Speedup (p) = Time (serial)/ Time (parallel)

Correctness in parallelisation requires synchronisation (locking)

Synchronisation and atomic operations causes loss of performance, communication latency.

Amdahl's law, establishes the maximum improvement to a system when only part of the system has been improved. Gustafson and Barsis noted that Amdahl's Law assumed a computation problem of fixed data set size.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

A detailed review was conducted in 2016 looking at the infrastructure of the Melbourne Research Cloud, High Performance Computing, and Research Data Storage Services. University desired a 'more unified experience to access compute services'

Recommended solution, based on technology and usage, is to make use of existing cloud with an expansion of general provisioning and use of a small number of bare metal nodes.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The 'bare metal' HPC component really will be laconic. "Real" HPC is a mere c276 cores, 21 GB per core. 2 socket Intel E5-2643 v3 CPU with 6-core per socket, 192GB memory, 2x 1.2TB SAS drives, 2x 40GbE network. "Cloud" partitions is almost 400 virtual machines with over 3,000 cores. There is also a GPU partition (big expansion this year), and departmental partitions (water and ashley).

This is not a big cluster by international standards! c.f., *The Provision of HPC Resources to Top Universities* <http://levlafayette.com/node/528>

But it is important! Spartan and the model of an HPC-Cloud Hybrid has been featured at Multicore World, Wellington, 2016, 2017; eResearchAustralasia 2016, Center for Scientific Computing (CSC) Goethe University Frankfurt, 2016, High Performance Computing Center (HLRS) University of Stuttgart, 2016, High Performance Computing Centre Albert-Ludwigs-University Freiburg, 2016; Eur Research (CERN), 2016, Centre Informatique National de l'Enseignement Supérieur, 2016, Centro Nacional de Supercomputación, Barcelona, 2016.

▪ <https://www.youtube.com/watch?v=6D1lobnGZqE>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](#)

Also featured in OpenStack and HPC Workload Management in Stig Telfer (ed), *The Crossroads of Cloud and HPC: OpenStack for Scientific Research*, Open Stack, 2016
<http://openstack.org/assets/science/OpenStack-CloudandHPC6x9Booklet-v4-online.pdf>

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Service	Network Device	Network	Protocol	Latency (usecs)
UoM HPC Traditional	Mellanox	56Gb	Infiniband FDR	1.17
Legacy Edward HPC	Cisco Nexus	10Gbe	TCP/IP	19
Spartan Cloud nodes	Cisco Nexus	10Gbe	TCP/IP	60
Spartan Bare Metal	Mellanox	40Gbe	TCP/IP	6.85
Spartan Bare Metal	Mellanox		hernet	1.84
Spartan Bare Metal	Mellanox		hernet	1.15
Spartan Bare Metal	Mellanox	56Gbe	net	1.68
Spartan Bare Metal	Mellanox	100Gbe	net	1.3

Job	Task	Resources	Control	HPC	Spartan Cloud
BWA	Disk	8 core Single Node	1:18:49	1:02:56	1:40:21
GROMACS	Compute	128 core Multinode	0:30:02	0:30:10	0:30:32
NAMD	Compute, I/O	128 core Multinode	1:11:41	1:00:46	1:55:54

Spartan (like Edward) uses its own authentication that is tied to the university Security Assertion Markup Language (SAML). The login URL is `https://dashboard.hpc.unimelb.edu.au/karaage`

Assignment Project Exam Help

Users on Spartan must belong to a project. Projects must be led by a University of Melbourne researcher (the "Principal Investigator") and approved by the Head of Research Compute Services.

<https://eduassistpro.github.io/>

Participants in a project can be researchers or researchers off from anywhere.

Add WeChat [edu_assist_pro](#)

The University, through Research Platforms, has an extensive training programme for researchers who wish to use Spartan. This includes day-long courses in “Introduction to Linux and HPC Using Spartan”, “Edward to Spartan Transition Workshop”, “Linux Shell Scripting for High Performance Computing”, and “Parallel Programming On Spartan”.

To log on to a HPC system, you will need a user account and password and a Secure Shell (ssh) client. Most HPC cluster administrators do not allow connections with protocols such as Telnet, FTP or RSH as they insecurely send passwords in plain-text over the network, which is easily captured by packet analyser tools (e.g., Wireshark). Linux distributions almost always include SSH as part of the default installation as does Mac OS 10.x, although you may also wish to use the Fugu SSH client. For MS-Windows users the free PuTTY client is recommended. To transfer files use scp, WinSCP, Filezilla, and especially rsync.

Assignment Project Exam Help

Logins to Spartan are based on <https://eduassistpro.github.io/>

`ssh your-username@spartan2.hpc.unimelb.edu.au` Add WeChat [edu_assist_pro](#)

Note `spartan2`. This is a second login node that was created specifically for this class.

For help go to <http://dashboard.hpc.unimelb.edu.au> or check `man spartan`. Lots of example scripts at `/usr/local/common`

Need more help? Problems with submitting a job, need a new application or extension to an existing application installed, if job generated unexpected errors etc., an email can be sent to: `hpc-support@unimelb.edu.au`

Assignment Project Exam Help

Assumption here is that everyone has had exposure to the Linux command line. If not, you'd better get some! At least learn the twenty or so basic environment commands to navigate the environment, manipulate files, manage processes. Plenty of good online material available (e.g., my book "Supercomputing with Linux", <https://github.com/VPAC/superlinux>)

<https://eduassistpro.github.io/>

Environment modules provide the user's environment (e.g., paths) via module files. Each module file contains configuration information for the user's session to operate according to the application's needs, such as the location of the application's executables, its manual path, the library path, and so forth.

Add WeChat [edu_assist_pro](https://eduassistpro.github.io/)

Modulefiles also have the advantages of being shared with many users on a system and easily allowing multiple installations of the same application but with different versions and compilation options. Sometimes users want the latest and greatest of a particular version of an application for the feature-set they offer. In other cases, such as someone who is participating in a research project, a consistent version of an application is desired. Having multiple version of applications available on a system is essential in research computing.

Some basic module commands include the following:

`module help`

The command `module help`, by itself, provides a list of the switches, subcommands, and subcommand arguments that are available through the environment modules package.

`module avail`

This option lists all the modules

<https://eduassistpro.github.io/>

d.

`module whatis <modulefile>`

This option provides a description of the module listed.

Add WeChat edu_assist_pro

`module display <modulefile>`

Use this command to see exactly what a given modulefile will do to your environment, such as what will be added to the PATH, MANPATH, etc. environment variables.

`module load <modulefile>`

This adds one or more modulefiles to the user's current environment (some modulefiles load other modulefiles).

`module unload <modulefile>`

This removes any listed modules from the user's current environment.

`module switch <modulefile1> <modulefile2>`

This unloads one modulefile (modulefile1) and loads another (modulefile2).

`module purge`

This removes all modules from the user's environment.

In the lmod system as used on Spartan there is also “module spider” which will search for all possible modules and not just those in the existing module path.

(Image from NASA, Apollo 9 “spider module”)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

The Portable Batch System (or simply PBS) is a utility software that performs job scheduling by assigning unattended background tasks expressed as batch jobs among the available resources. The scheduler provides for parameterisation of computer resources, an automatic submission of execution tasks, and a notification system for incidents.

The original Portable Batch System was developed by MRI Technology Solutions under contract to NASA in the early 1990s. In 1998 the original version of PBS was released as an open-source product as OpenPBS. This was formally, Cluster Resources) who developed TORQUE (Terascale Queue Manager). Many of the original engineering team and what commercial products from the original product is now part of Altair Engineering who have their own. In addition to this the popular job scheduler Slurm (originally “Simple Linux Utility for Resource Management”), now simply called Slurm Workload Manager, also uses batch script where are very similar in intent and style to PBS scripts.

Spartan uses the Slurm Workload Manager. A job script written on one needs to be translated to another (handy script available pbs2slurm <https://github.com/bjpop/pbs2slurm>)

In addition to this variety of implementations of PBS different institutions may also make further elaborations and specifications to their submission filters (e.g., site-specific queues, user projects for accounting). (Image from the otherwise dry IBM 'Red Book' on Queue Management)

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Submitting and running jobs is a relatively straight-forward process consisting of:

- 1) Setup and launch
- 2) Job Control, Monitor results
- 3) Retrieve results and analyse.

Don't run jobs on the login node! Use the queuing system to submit jobs.

1. Setup and launch consists of `qsub` which makes resource requests and then commands, and optionally `qrun` for running the job.

Core command for checking queue: `qstat`
Alternative command for checking queue: `showq`
Core command for job submission: `sbatch [jobscript]`

2. Check job status (by ID or user), cancel job.

Core command for checking job in Slurm: `squeue -j [jobid]`
Detailed command in Slurm: `scontrol show job [jobid]`
Core command for deleting job in Slurm: `scancel [jobid]`

3. Slurm provides an error and output files. They may also have files for post-job processing. Graphic visualisation is best done on the desktop.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

then runs the executable

the job from. If your data is a

Modifying resource allocation requests can improve job efficiency.

For example shared-memory multithreaded jobs on Spartan (e.g., OpenMP), modify the `--cpus-per-task` to a maximum of 8, which is the maximum number of cores on a single instance.

```
#SBATCH --cpus-per-task=8
```

For distributed-memory multicore job using message passing, the multinode partition has to be invoked and the resource request

```
#!/bin/bash
#SBATCH -p physical
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=4
module load my-app-compiler/version
srun my-mpi-app
```

Note that multithreaded jobs *cannot* be used in a distributed memory model across nodes. They can however exist be conducted on distributed memory jobs which include a shared memory component (hybrid OpenMP-MPI jobs).

Assignment Project Exam Help
<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

Alternative job submissions include specifying batch arrays, and batch dependencies.

In the first case, the same batch script, and therefore the same resource requests, is used multiple times. A typical example is to apply the same task across multiple datasets. The following example submits 10 batch jobs with myapp running against datasets dataset1.csv, dataset2.csv, ... dataset10.csv

```
#SBATCH --array=1-10
myapp ${SLURM_ARRAY_TASK_ID}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

In the second case a dependency condition is established, creating a conditional pipeline. The dependency conditions consist of `after`, `afterok`, `afternotok`, `before`, `beforeok`, `beforenotok`. A task is required as the input of the next job.

ch the launching of a batch script
as consist of `after`, `afterok`,
e is where the output of one job

```
#SBATCH --dependency=afterok:myfirstjobid mysecondjob
```

For real-time interaction, with resource requests made on the command line, an interactive job is called. This puts the user on to a compute node.

This is typically done if they user wants to run a large script (and shouldn't do it on the login node), or wants to test or debug a job. The following command would launch one node with two processors for ten minutes.

```
[lev@spartan interact]$ srun --ntasks-per-node=2  
srun: job 164 queued an  
srun: job 164 has been  
[lev@spartan-rc002 interact]$
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

In almost all cases it is much better to do computation on the cluster and visualisation on a local system. In some cases however it is unavoidable to require x-windows forwarding.

It is best to login with the -Y option for security and then to login with -X to the login node. The computer node with then pass through the graphics via the login node to the desktop system.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Please note that you will need an x-windows client your desktop for the visualisation.

Add WeChat edu_assist_pro

```
[lev@cricetomys ~]$ ssh  
lev@spartan.hpc.unimelb.edu.au -Y
```

```
[lev@spartan]$ sinteractive --nodes=1 --ntasks-per-  
node=2 --x11=first  
srun: job 602795 queued and waiting for resources  
srun: job 602795 has been allocated resources  
[lev@spartan-rc002 ~]$ xclock
```

User Commands

Job submission

PBS/Torque

qsub [script_file]

SLURM

sbatch [script_file]

Job submission

qdel [job_id]

scancel [job_id]

Job status (by job)

qstat [job_id]

squeue [job_id]

Job status (by user)

qstat -u [user_name]

squeue -u [user_name]

Node list

pbsn

Queue list

qstat -Q

Cluster status

showq, qstatus -a

artition]

Environment

Job ID

\$PBS_JOBID

\$SLURM_JOBID

Submit Directory

\$PBS_O_WORKDIR

\$SLURM_SUBMIT_DIR

Submit Host

\$PBS_O_HOST

\$SLURM_SUBMIT_HOST

Node List

\$PBS_NODEFILE

\$SLURM_JOB_NODELIST

Job Array Index

\$PBS_ARRAYID

\$SLURM_ARRAY_TASK_ID

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Job Specification

PBS

SLURM

Script directive

#PBS

#SBATCH

Queue

-q [queue]

-p [queue]

Job Name

-N [name]

--job-name=[name]

Nodes

-l nodes=[count]

-N [min]-[max]

CPU Count

-l pp

Wall Clock Limit

-l walltime=[hh:mm:ss] m:ss]

Event Address

-M [address] [address]

Event Notification

-m abe

--mail-type=[events]

Memory Size

-l mem=[MB]

--mem=[mem][M|G|T]

Proc Memory Size

-l pmem=[MB]

--mem-per-cpu=[mem][M|G|T]

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

One form of parallel programming is multithreading, whereby a master thread forks a number of sub-threads and divides tasks between them. The threads will then run concurrently and are then joined at a subsequent point to resume normal serial application.

One implementation of multithreading is OpenMP (Open Multi-Processing). It is an Application Program Interface that includes directives for multi-threaded, shared memory parallel programming. The directives are included in the C or Fortran source code and in a system where OpenMP is not implemented, they would be interpreted as comments.

There is no doubt that OpenMP programming, however it is limited to a single system unit (no distributed memory) and d rather than using message passing. Many examples in ``/usr/local/comp`

(image from: User A1, Wikipedia)

```
#include <stdio.h>
#include "omp.h"
int main(void)
{
    int id;
    #pragma omp parallel num_threads(8) private(id)
    {
        int id = omp_get_thread_num();
        printf("Hello world
    }
return 0;
}
```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
program hello2omp
    include "omp_lib.h"
    integer :: id
    !$omp parallel num_threads(8) private(id)
        id = omp_get_thread_num()
        print *, "Hello world", id
    !$omp end parallel
end program hello2omp
```

Moving from shared memory to parallel programming involves a conceptual change from multi-threaded programming to a message passing paradigm. In this case, MPI (Message Passing Interface) is one of the most well popular standards and is used here, along with a popular implementation as OpenMPI.

The core principle is that many processes should be able cooperate to solve a problem by passing messages to each through a common co

Assignment Project Exam Help

<https://eduassistpro.github.io/>

The flexible architecture does overcome serial bottlenecks, but it also does require explicit programmer effort (the "questing beast" of automatic parallelisation remains somewhat elusive).

Add WeChat edu_assist_pro

The programmer is responsible for identifying opportunities for parallelism and implementing algorithms for parallelisation using MPI.


```

#include <stdio.h>
#include "mpi.h"
int main( argc, argv )
int  argc;
char **argv;
{
    int rank, size;
    MPI_Init( &argc, &argv );
    MPI_Comm_size( MPI_COMM_
    MPI_Comm_rank( MPI_COMM_
    printf( "Hello world from process
    MPI_Finalize();
    return 0;
}

```

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```

! Fortran MPI Hello World
program hello
include 'mpif.h'
integer rank, size, ierror, tag, status(MPI_STATUS_SIZE)

call MPI_INIT(ierror)
call MPI_COMM_SIZE(MPI_COMM_WORLD, size, ierror)
call MPI_COMM_RANK(MPI_COMM_WORLD, rank, ierror)
print*, 'node', rank, ': Hello world'
call MPI_FINALIZE(ierror)
end

```

The OpenMP example needs to be compiled with OpenMP directives. The OpenMP example cannot run across compute nodes; therefore it is best run on the “cloud” partition. The OpenMPI compilation needs to call the MPI wrappers.

```
module load OpenMPI/1.10.0-GCC-4.9.2
gcc -fopenmp helloomp.c -o helloomp
mpigcc mpihelloworld.c -o mpihelloworld
```

Assignment Project Exam Help

```
#!/bin/bash
#SBATCH -p cloud
#SBATCH --nodes=1
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=16
export OMP_NUM_THREADS=16
module load GCC/4.9.2
mpiexecu helloomp
```

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

```
#!/bin/bash
#SBATCH -p physical
#SBATCH --nodes=2
#SBATCH --ntasks=16
module load OpenMPI/1.10.2-GCC-4.9.2
mpiexec mpi-helloworld
```

Java can be compiled with MPI bindings; we have done this with OpenMPI/3.0.0 only. A more common option is to use MPJ-Express. The following “HelloWorld.java” program is compiled and executed.

```
import mpi.*;
public class HelloWorld {
public static void main(String args[]) throws Exception {
    MPI.Init(args);
    int me = MPI.C
    int size = MPI
    System.out.println("Hi from
    MPI.Finalize());
}
}
```

```
sinteractive --time=1:00:00 --nodes=1 --ntasks=2
module load MPJ-Express
javac -cp ./usr/local/easybuild/software/MPJ-Express/0.44-goolf-
2015a-Java-9.0.4/lib/mpj.jar HelloWorld.java
mpjrun.sh -np 2 HelloWorld
```

Python too has various MPI bindings available. The most common used is MPI4Py. Sample assignment data is provided from the 2016 and 2015 in the Spartan directory, ``/usr/local/common``. As a package it can be simply imported (e.g., ``from mpi4py import MPI``).

But remember! With environment modules with extensions you do not necessarily get all the packages/libraries/extensions that you might expect. See the README file for an explanation of how to review the extensions already installed.

Examples provided in the directory <https://eduassistpro.github.io/>
Python jobs with MPI bindings with Slurm
submissions scripts for single-core, dual-core, four-core, eight-core, sixteen-core (different partition), and two-nodes with four cores each.

The actual Python script for the above is ``twitter_search_541635.py``

A very popular and basic use of MPI Send and Recv routines is a ping-ping program. Why? Because it can be used to test latency within and between nodes and partitions if they have different interconnect (like on Spartan).

An example is given in `/usr/local/common/MPI` as `mpi-pingpong.c` with a job submission script that can be modified accord to the test case `mpi-pingpong.slurm`. There are some routines here which manage the communication in the ping-pong activity.

`MPI_Status()` `MPI_Status` is structure and is typically attached to an `MPI_Recv()` routine.

`MPI_Request()` A wrapper for `MPI_Request` is any, waitall, waitsome, start, cancel, startall.



`MPI_Barrier()` Enforces synchronisation between MPI processes in a group by placing a barrier on communication between groups. An MPI barrier completes after all group members have entered the barrier.

`MPI_Wtime()` - Returns an elapsed time as a floating-point number of seconds on the calling processor from an arbitrary time in the past.

There are *many* other MPI routines which are *not* going to be explored here! However just as a little taste one of the most popular is MPI collective communications and reduction operations. Collective communications include MPI_Broadcast, MPI_Scatter, MPI_Gather, MPI_Reduce, and MPI_Allreduce.

MPI_Bcast Broadcasts a message from the process with rank "root" to all other processes of the communicator, including itself. It is significantly more preferable than using a loop.

MPI_Scatter sends data from MPI_Gather. The outcome is the inverse operation of MPI_Gather. The outcome is that each process in a group has a unique piece of the data. MPI_Scatterv scatters a buffer into all tasks in a group.




MPI_Reduce performs a reduce operation (such as sum, max, logical AND, etc.) across all the members of a communication group.

MPI_Allreduce conducts the same operation but returns the reduced result to all processors.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



The general principle in Reduce and All Reduce is the idea of reducing a set of numbers to a small set via a function. If you have a set of numbers (e.g., [1,2,3,4,5]) a reduce function (e.g., sum) can convert that set to a reduced set (e.g., 15). MPI_Reduce takes in an array of values as that set and outputs the result to the root process. MPI_AllReduce outputs the result to all processes.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro